

Advanced Keras: motivation

In introducing Deep Learning, we have asserted that

It's all about the Loss function

That is: the key to solving many Deep Learning problems

- Is not in devising a complex network architecture
- But in writing a Loss function that captures the semantics of the problem

Up until now

- We have been using pre-defined Loss functions (e.g., `binary_crossentropy`)
- Specifying the Loss function in the compile statement

```
model.compile(loss='binary_crossentropy')
```

- Using the built-in "training loop" (cycling through epochs, using Gradient Descent to update weights)
 - `model.fit(...)`

You can [write your own loss functions](https://keras.io/api/losses/) (<https://keras.io/api/losses/>).

In Keras, a Loss function has the signature

loss_fn(y_true, y_pred, sample_weight=None)

But what if your Loss function needs access to values that are not part of the signature ?

In that case, you might need to write your own training loop

- Cycle through epochs
- Within each epoch, cycle through mini-batches of examples
- For each mini-batch of examples: compute the custom loss
- Compute the gradient of the loss with respect to the weights
- Update the weights

In this module we will

- Illustrate the Functional model
 - allows more complex network architectures than the Keras Sequential model
- Show how to write complex Loss functions
- Show how to write custom training loops

In addition, we will show you how to create new Layer types (e.g., Transformer)

Functional model

[Autoencoder \(Autoencoder_example.ipynb#Library-for-Vanilla-Autoencoder-demo\)](#)

- Functional model

Issues

- We could use a Sequential model with initial Encoder layers and final Decoder layers
 - But we would not be able to independently access the Encoder nor the Decoder as isolated models

VAE: Complex Loss; Manual Gradient updates

Variational Autoencoder (VAE) (<https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/vae.ipynb#scrollTo=DEU05Oe0vJrY>).

- Functional model
- VAE: Custom train step (<https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/vae.ipynb#scrollTo=0EHkZ1WCHw9E>).
 - Complex loss

Issues

Transformer: Custom layers, Skip connections, Layer Norm

[Transformer layer \(https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/nlp/ipynb/neural_machine_translation_with_transformer.ipynb#scrollto=IMkSs\)](https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/nlp/ipynb/neural_machine_translation_with_transformer.ipynb#scrollto=IMkSs)

- Functional model
- Custom layers
- Layer Norm
- Skip connections

The following diagram shows the architecture, which we can compare to the code

- [Full architecture diagram \(compare with code\) \(Transformer.ipynb#Full-Encoder-Decoder-Transformer-architecture\)](#)

We can dig deeper to examine how the Attention layers are implemented in code:

- [Scaled dot-product attention \(https://www.tensorflow.org/text/tutorials/transformer#scaled_dot_product_attention\)](https://www.tensorflow.org/text/tutorials/transformer#scaled_dot_product_attention)
- [Multi-head attention \(https://www.tensorflow.org/text/tutorials/transformer#multi-head_attention\)](https://www.tensorflow.org/text/tutorials/transformer#multi-head_attention)

Issues

- Build a new layer type
 - Why are the components layers (e.g., Dense, MultiHeadAttention, LayerNormalization) instantiated in the class constructor
 - As opposed to being defined in the "call" method
 - Because we **need** one instance of the layer
 - Not a new instance each time the class is "called" per batch
 - This would result in brand new weights for each example batch
 - The "call" method accesses the shared layer instances and performs the computation using them
-

Neural Style Transfer

Neural Style Transfer (https://keras.io/examples/generative/neural_style_transfer/)

- Complex Loss
- Custom training loop

Here (https://www.tensorflow.org/tutorials/generative/style_transfer) is a tutorial view of the notebook.

In [2]: `print("Done")`

Done