

# Task agnostic models

Early approaches to NLP via Deep Learning created task-specific architectures.

The power of these models was enhanced by increasingly sophisticated word representations

- Obtained via a Language Model

As Language Models have grown increasingly powerful (and large)

- The realization is that the architecture for the Language Model is universal !
  - No need to augment the Language Model with a *deep* task-specific "Head"
- Just use Transfer Learning on the Language Model !

This approach is called *Supervised Pre-training + Fine-Tuning*

- *Supervised Pre-Training*: the Language Model (e.g., predict the next word)
- *Fine-Tuning*: add a task specific head and fine-tune

Contrast this to Word Embeddings, which also use Transfer Learning

- Embeddings transfer *word-level* concepts
- Transferring entire Language Models transfer *semantic* concepts

Because the Pre-Trained model has a very specific input format (and output)

- You often have to encode your task-specific input to fit

For example:

- Consider a Pre-Trained model that performs text completion (predict the next)
- Turn your task into a text completion problem
- [See \(<https://arxiv.org/pdf/2005.14165.pdf>\)](https://arxiv.org/pdf/2005.14165.pdf) Appendix G (pages 75+) for examples

## Task: Unscramble the letters

---

Context:	Please unscramble the letters in the word and write that word
----------	---

---

skicts =

---

Target completion: sticks

## Task: English to French

Context:	English: Please unscramble the letters in the word and write that word
	French:
Target completion:	Veuillez déchiffrer les lettres du mot et écrire ce mot

Sometimes the task encodings are not completely obvious (see [GPT Section 3.3 \(https://cdn.openai.com/research-covers/language-unsupervised/language\\_understanding\\_paper.pdf\)](https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf))

- Task: Are two sentences similar ?
  - Issue
    - There is no natural ordering of the two sentences
    - So concatenating the two (with a delimiter) is misleading
  - Solution
    - Obtain two representations of the sentence pair, once for each ordering
    - Add them together element-wise
    - Feed sum into Classifier

- Task: multiple choice questions answering: given context, question plus list of possible answers
  - Solution:
    - Obtain representation for each answer
      - Concatenate (with delimiter): context, questions, answer
    - Feed each representation into a softmax to obtain probability distribution over answers



## GPT: Task encoding

---

---

Picture from: [https://cdn.openai.com/research-covers/language-unsupervised/language\\_understanding\\_paper.pdf](https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf)

---

Universal model: adapting task-specific inputs

---

Picture from: <http://jalammar.github.io/images/bert-tasks.png>

## **From a very practical standpoint**

- In the near future (maybe even now) you will not create a new model
- You will use an existing Language Model
  - Trained with lots of data
  - At great cost
- And fine-tune to your task

# **Models using Supervised Pre-training + Fine-Tuning**

We present a few models using this approach.

# GPT: Generalized Pre-Training

GPT is a sequence of increasingly powerful (and big) models of similar architecture.

- The Decoder side of a Transformer Encoder-Decoder model
  - Masked Self-attention
  - Left to Right, unidirectional

Each generation

- Increase the number of Transformer blocks
- Increases the size of the training data

All models use

- Byte Pair Encoding
- Initial encode words with word embeddings

They are all trained on a Language Model objective.

## GPT: architecture

---

Picture from: [https://cdn.openai.com/research-covers/language-unsupervised/language\\_understanding\\_paper.pdf](https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf)





The models can be described as

$$\begin{aligned}h_0 &= UW_e + W_p \\h_i &= \text{transformer\_block}(h_i) \quad \text{for } 1 \leq i \leq n \\ \text{prod}(U) &= \text{softmax}(h_n W_e^T)\end{aligned}$$

where

$$\begin{aligned}U &\text{ context of size } k : [u_{-k}, \dots, u_{-1}] \\h_i &\text{ Output of transformer block } i \\n &\text{ number of transformer blocks/layers} \\W_e &\text{ token embedding matrix} \\W_p &\text{ position encoding matrix}\end{aligned}$$

Let's understand this

- $h_0$ , the output of the input layer
  - Uses word embeddings  $W_e$  on the input  $U$
  - Adds *positional* encoding  $W_p$  to the tokens
- There are layers  $h_i$  of Transformer blocks  $1 \leq i \leq n$
- The output  $\text{prod}(U)$ 
  - Takes the final layer output  $h_n$
  - Reverses the embedding  $W_e^T$  to get back to original tokens
  - Uses a softmax to get a probability distribution over the tokens  $U$ 
    - Distribution over the predicted next token

The training objective is to maximize log likelihood on  $\mathcal{U}$

$$\mathcal{L}_1(\mathcal{U}) = \sum_i \log P(u_i | u_{i-k}, \dots, u_{i-1}; \Theta)$$

[paper \(https://cdn.openai.com/research-covers/language-unsupervised/language\\_understanding\\_paper.pdf\)](https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf).

[Summary \(https://openai.com/blog/language-unsupervised/\)](https://openai.com/blog/language-unsupervised/).

- 12 Transformer blocks (37 layers)
  - $n_{\text{heads}} = 12, d_{\text{head}} = 64$ 
    - $d_{\text{model}} = n_{\text{heads}} * d_{\text{head}} = 768$
    - $d_{\text{model}}$  is size of bottle-neck layer
- 117 million weights
- Trained on
  - 5GB of text (BooksCorpus dataset consisting of 7,000 books)
  - Sequence of 512 tokens
  - Training time
    - 30 days on 8 GPUs
    - 26 petaflop-days

The original Unsupervised Training to create the Language Model.

This is followed by Fine Tuning on a smaller task-specific training set  $\mathcal{C}$

This can be described as:

- Add linear output layer  $W_y$  to the model used for Language Modeling:
- $h_l^m$  is output of transformer block  $l$  on input of length  $m$
- Using  $\Theta$  from unsupervised pre-training
- Fine Tuning Objective:

- maximize log likelihood on  $\mathcal{C}$

$$\mathcal{L}_2(\mathcal{C}) = \sum_{(x,y)} P(y|x_1, \dots, x_m) = \text{softmax}(h_l^m W_y)$$

The authors also experimented with a Fine Tuning Objective that included the Language Model

$$\mathcal{L}_3(\mathcal{C}) = \mathcal{L}_2(\mathcal{C}) + \lambda \mathcal{L}_1(\mathcal{C})$$

## **Results of Supervised Pre-Training + Fine-Tuning**

- Tested on 12 tasks
- Improved state-of-the-art results on 9 out of the 12

# BERT

[paper \(https://arxiv.org/pdf/1810.04805.pdf\)](https://arxiv.org/pdf/1810.04805.pdf)

BERT (Bidirectional Encoder Representations from Transformers) is also a *fine-tuning* (universal model) approach, like GPT

- does not use *masked attention* to force causal ordering
- uses a Masked Language Model pre-training objective

The Transformer in OpenAI's GPT uses *Masked Self-Attention*

- the Language Models/training objectives are conditioned on *prefix* and *suffix*, not full context
- So is fundamentally a left-to-right Language Model



## Masked Language Model task

- Mask (obscure) 15% of the input tokens, chosen at random
- The method for masking takes one of three forms
  - 80% of the time, hide it: replace with [MASK] token
  - 10% of the time: replace it with a random word
  - 10% of the time: don't obscure it

The training objective is to predict the masked word

The authors explain

- Since encoder does not know which words have been masked
- Or which of the masked words were random replacements
- It must maintain a context for **all** tokens

They also state that, since random replacement only occurs 1.5% of the time ( $10\% * 15\%$ ), this does not seem to destroy language understanding

## BERT in action

Interactive model for MLM (<https://huggingface.co/bert-base-uncased?text=Washington+is+the+%5BMASK%5D+of+the+US>).